Compiling PDDL3 Qualitative Preferences without Using Automata

...

Abstract

We address the problem of planning with preferences in propositional domains extended with the class of (preferred) temporally extended goals supported in PDDL3, that is part of the standard planning language PDDL since the 5th International Planning Competition (IPC5). Such preferences are useful to characterise plan quality by allowing the user to express certain soft constraints on the state trajectory of the desired solution plans. Starting from the work of Keyder and Geffner on compiling (reachability) soft goals, we propose a compilation scheme for translating a STRIPS problem enriched with qualitative PDDL3 preferences (and possibly also with soft goals) into an equivalent STRIPS problem with action costs. The proposed compilation, which supports all types of preferences in the benchmarks from IPC5, allows many existing STRIPS planners to immediately address planning with preferences. An experimental analysis presented in the paper evaluates the performance of state-of-the-art STRIPS planners supporting action costs using our compilation approach to deal with qualitative PDDL3 preferences. The results indicate that our approach is highly competitive with respect to current planners that natively support the considered class of preferences.

Introduction

Planning with preferences, also called "over-subscription planning" in [2, 5?], concerns the generation of plans for problems involving soft goals or soft state-trajectory constraints (called preferences in PDDL3), that it is desired a plan satisfies, but that do not have to be satisfied. The quality of a solution plan for these problems depends on the soft goals and preferences that are satisfied.

For instance, a useful class of preferences than can be expressed in PDDL3 [6] consists of *always preferences*, requiring that a certain condition should hold in *every* state reached by a plan. As discussed in [??6?], adding always preferences to the problem model can be very useful to express safety or maintenance conditions, and other desired plan properties. An simple example of such conditions is "whenever a building surveillance robot is outside a room, all the room doors should be closed".

PDDL3 supports other useful types of preferences, and in particular the qualitative preferences of types *at-end*, which

are which are equivalent to soft goals, *sometime*, *sometime-before* and *at-most-once*, which are all the types used in the available benchmarks for planning with qualitative PDDL3 preferences [6]. Examples of preferences that can be expressed through these constructs in a logistics domain are: "sometime during the plan the fuel in the tank of every vehicle should be full", "a certain depots should be visited before another once", "every store should be visited at most once" (the reader can find additional examples in [6]).

In this paper, we study propositional planning with these types of preferences through a compilation approach.

Related Work

Our compilative approach is inspired by the work of Keyder and Geffner [?] on compiling soft goals into STRIPS with action costs (here denoted with STRIPS+). In this work the compilation scheme introduces, for each soft goal p of the problem, a dummy goal p' that can be achieved using two actions in mutual exclusion. The first one, which is called *collect(p)*, has cost equal to 0 and requires that p be true when it is applied; the second one, whic is called *forgo(p)*, has cost equal to the utlity of p and requires that p be false when it is applied.

Both of these action can be performed at the end of the plan and for each soft goal p but just one of {*collect(p), forgo(p)*} can appear in the plan depending on whether the soft goal has been achived or not. This scheme has achieved good perfomance which can be improved with the use of an ad hoc admissible heuristic based on the reachability of soft goals [9].

The most prominent existing planners supporting PDDL3 preferences are HPlan-P [? ?], which won the "qualitative preference" track of IPC-5, MIPS-XXL [? ?] and the more recent LPRPG-P [?] and its extension in [?]. These (forward) planners represent preferences through automata whose states are synchronised with the states generated by the action plans, so that an accepting automaton state corresponds to preference satisfaction. For the synchronisation, HPlan-P and LPRPG-P use planner-specific techniques, while MIPS-XXL compiles the automata by modifying the domain operators and adding new ones modelling the automata transitions of the grounded preferences.

Our computation method is very different from the one of MIPS-XXL since, rather than translating automata into new

operators, the problem preferences are compiled by only modifying the domain operators, possibly creating multiple variants of them. Moreover, our compiled files only use STRIPS+, while MIPS-XXL also uses numerical fluents.[1]

The works on compiling LTL goal formulas by Cresswell and Coddington [**?**] and Rintanen [**?**] are also somewhat related to ours, but with important differences. Their methods handle *hard* temporally extended goals instead of preferences, i.e., every temporally extended goal must be satisfied in a valid plan, and hence there is no notion of plan quality referred to the amount of satisfied preferences. Rintanen's compilation considers only single literals in the always formulae (while we deal with arbitrary CNF formulas), and it appears that extending it to handle more general formulas requires substantial new techniques [**?**]. An implementation of Crosswell and Coddington's approach is unavailable, but Bayer and McIlaraith [**?**] observed that their approach suffers exponential blow up problems and performs less efficiently than the approach of HPlan-P.

==IMPORTANTE, CONFRONTO LAVORO NEBEL:== An important work about soft-trajectory constraints compilation, which is closely related to ours, has been recently proposed by Wright, Mattüller and Nebel in [**?**] (cerca riferimento articolo Nebel). This approach is a based on the compilation of the soft trajectory constraints into conditional effects and state dependent action costs using LTL_f and Büchi automata. There are some similarities between between our and their approach but at the same time there are also differences.

FP: differenze approccio Nebel e nostro: 1) numero di fluenti introdotto nella compilazione diverso; 2) il nostro e' un approccio piu' specifico PDDL3-oriented, il loro piu' generale; 3) diverso meccanismo di aggiornamento dei costi, nel loro caso ricompensa e penalita', nel nostro solo penalit 4) il loro schema prevede quindi costi negativi e postivi, il nostro positivi; nel loro caso per avere solo costi positivi e' necessario perdere l'ottimalita' nel nostro caso no ed inoltre volendo potremmo introdurre dei costi negativi per quelle preferenze la cui violazione puo' essere testata solo alla fine del piano mantenendo l'ottimalita'; ho fatto inoltre menzione del fatto che i costi possono essere incrementanti il prima possibile (es. always)

1) In our approach, given a preference P, we introduce at most a pair of boolean fluents, typically one additional fluent to represent if a preference is violated or not (P-*violated*) and in some cases an additional fluent to correctly represent the status of a preference during the planning, while in their approach it is introduced a boolean variable for each state of the corresponding automaton of P. **2)** Their approach is more general while ours is focused on PDDL3 constraints and therefore it is more specific.

3) In their work the cost of the plan during the planning is updated by using rewards and penalties (negative and positive costs respectively). The cost of the plan is increased whenever a violation of a preference (also reversible) occurs. On the contrary, if an operator causes the satisfaction of a preference then the cost of the plan is decreased. In both cases the negative or positive cost is equal to the utility of the interested preference[2]. **4)** This type of cost update requires the use of negative costs while our compilation scheme produces a problem whose costs are monotonically increasing because, similarly to what was proposed in Keyder and Geffner, costs are realized only at the end of the planning. In our scheme some costs can be anticipated for those preferences whose violation is irreversible (e.g. always, sometime-before) and negative costs could be used for those preferences that may be "temporarily" violated (e.g. sometime, at-end). In both cases our scheme would maintain the optimality but in this work we have provided a compilation based only on positive costs in order to take advantage of a wider spectrum of classical planners (few planners still support negative costs).

we propose a compilation scheme for translating a STRIPS problem with PDDL3 qualitative preferences into an equivalent STRIPS+ problem. Handling action costs is a practically important, basic functionality that is supported by many powerful planners; the proposed compilation method allows them to immediately support (through the compiled problems) the considered class of preferences with no change to their algorithms and code.

Propositional Planning with Qualitative PDDL3 Preferences

A STRIPS problem is a tuple $\langle F, I, O, G \rangle$ where F is a set of fluents, $I \subseteq F$ and $G \subseteq F$ are the initial state and goal set, respectively, and O is a set of actions or operators defined over F as follows.

A STRIPS operator $o \in O$ is a pair $\langle Pre(o), Eff(o) \rangle$, where $Pre(o)$ is a sets of atomic formulae over F and $Eff(o)$ is a set of literals over F. $Eff(o)^+$ denotes the set of positive literals in $Eff(o)$, $Eff(o)^-$ the set of negative literals in $Eff(o)$. An action sequence $\pi = \langle a_0, \ldots, a_m \rangle$ is applicable in a planning problem Π if all actions a_i are in O and there exists a sequence of states $\langle s_0, \ldots, s_{m+1} \rangle$ such that $s_0 = I$, $prec(a_i) \subseteq s_i$ and $s_{i+1} = s_i \setminus \{p \mid \neg p \in Eff(a_i)^-\} \cup Eff(a_i)^+$, for $i = 0 \ldots m$. Applicable action sequence π achieves a fluent g if $g \in s_{m+1}$, and is a valid plan for Π if it achieves each goal $g \in G$ (denoted with $\pi \models G$).

A STRIPS+ problem is a tuple $\langle F, I, O, G, c \rangle$, where $\langle F, I, O, G \rangle$ is a STRIPS problem and c is a function mapping each $o \in O$ to a non-negative real number. The cost $c(\pi)$ of a plan π is $\sum_{i=0}^{|\pi|-1} c(a_i)$, where $c(a_i)$ denotes the cost of the ith action a_i in π and $|\pi|$ is the length of π.

---

[1]Another compilation scheme using numerical fluents is considered in [6] to study the expressiveness of PDDL3(without an implementation).

[2]==Messo in footnote altrimenti era troppo lungo:== Note that both the violation and the satisfaction of a preference may be temporary conditions depending on the type of interested preference. For example an always preference can be irreversibly violated and its satisfaction can only be evaluated at the end of the plan considering the whole trajectory of the states; on the contrary a sometime-after preference can be satisfied and violated several times during the execution of the plan.

Starting from a STRIPS+ problem we define the notion of state trajectory generarated by a plan. Given a STRIPS+ $\Pi = \langle F, I, O, G, c \rangle$ problem, a plan π generates the trajectory $\langle (S_0, 0), (S_1, t_1), ..., (S_n, t_n) \rangle$ iif $S_0 = I$ and for each happening h generated by π, with h at time t, there is some i such that $t_i = t$ and S_i is the result of applying the happening h to S_{i-1}, and for every $j \in \{1 \ldots n\}$ there is a happening π at t_j.

The following is a fragment of the grammar of PDDL3, describing the new modalities of PDDL3 for expressing these soft state-trajectory constraints (indicated with con-GD) (the full BNF grammar is given in [? ?]) where <GD> is a goal description (a first order logic formula):

```
<con-GD> ::= (at end <GD>) | (always <GD>) |
             (sometime <GD>) |
             (at-most-once <GD>) |
             (sometime-after <GD> <GD>) |
             (sometime-before <GD> <GD>)
```

Let ϕ and ψ be atomic formulae over the predicates of a planning problem, then interpretation of the modal operators considered in this work is specified in Figure ?? [COMMENTATA]. We write $\pi \models P$ to indicate that the state-trajectory generated by π satisfies a preference P and we indicate with $\mathcal{A}, \mathcal{SB}, \mathcal{ST}, \mathcal{AO}, \mathcal{SG}$ the classes of all preferecense of always, sometime-before, sometime, at-most-once and soft goal respectively for a given STRIPS+ problem.

But in the following, without loss of generality, we will assume in the discussion that the formula ϕ and ψ involved in a preference P is expressed in conjuctive normal form $\phi = \phi_1 \wedge \phi_2 \wedge ... \wedge \phi_n$ where each ϕ_i ($i \in \{1 \ldots n\}$) is a clause of P formed by literals over the problem fluents.

Definition 1 *A* STRIPS+ *problem with preferences is a tuple* $\langle F, I, O, G, \mathscr{P}, c, u \rangle$ *where:*

- $\langle F, I, O, G, c \rangle$ *is a* STRIPS+ *problem;*

- $\mathscr{P} = \{\mathscr{P}_\mathcal{A} \cup \mathscr{P}_{\mathcal{SB}} \cup \mathscr{P}_{\mathcal{ST}} \cup \mathscr{P}_{\mathcal{AO}} \cup \mathscr{P}_{\mathcal{SG}}\}$ *is the set of the preferences of* Π *where* $\mathscr{P}_\mathcal{A} \subseteq \mathcal{A}$, $\mathscr{P}_{\mathcal{SB}} \subseteq \mathcal{SB}$, $\mathscr{P}_{\mathcal{ST}} \subseteq \mathcal{ST}$, $\mathscr{P}_{\mathcal{AO}} \subseteq \mathcal{AO}$ *and* $\mathscr{P}_{\mathcal{SG}} \subseteq \mathcal{SG}$;

- u *is an utility function mapping each* $P \in \mathscr{P}$ *to a value in* $\mathbb{R}_0^+$.

In the following the class of STRIPS+ problems with a set of preferences is indicated with STRIPS+P.

Definition 2 *Let* Π *be a* STRIPS+P *problem. The utility* $u(\pi)$ *of a plan* π *solving* Π *is the difference between the total amount of utility of the preferences by the plan and its cost:*

$$u(\pi) = \sum_{P \in \mathscr{P} : \pi \models P} u(P) - c(\pi).$$

The definition of plan utility for STRIPS+P is similar to the one given for STRIPS+ with soft goals by Keyder and Geffner [?]. A plan π with utility $u(\pi)$ for a STRIPS+P problem is optimal when there is no plan π' such that $u(\pi') > u(\pi)$. The definitions below are introduced to simplify the notation in the discussion.

Operator-Preference Interactions

The definitions below are introduced to simplify the notation in the discussion.

Definition 3 *Given a preference clause* $\phi_i = l_1 \vee l_2 \vee ... \vee l_m$, *the set* $L(\phi_i) = \{l_1, l_2, ..., l_m\}$ *is the equivalent set-based definition of* ϕ_i *and* $\overline{L}(\phi_i) = \{\neg l_1, \neg l_2, ..., \neg l_m\}$ *is the literal complement set of* $L(\phi_i)$.

Definition 4 *Given an operator* $o \in O$ *of a* STRIPS+P *problem,* $Z(o)$ *is the set of literal defined as:*

$$Z(o) = (Pre(o) \setminus \{p \mid \neg p \in \textit{Eff}(o)^-\}) \cup \textit{Eff}(o)^+ \cup \textit{Eff}(o)^-.$$

Note that the literals in $Z(o)$ hold in any reachable state resulting from the execution of operator o.

To shorten the following definitions and explanation we indicate the state where the operator o is applied with s and the state resulting from the application of o with s'.

Definition 5 *Given an operator* o *and CNF formula* $\phi = \phi_1 \wedge ... \wedge \phi_n$, *we define the set of clauses which will surely be true in the resulting state from the application of* o *as:*

$$C_\phi(o) = \{\phi_i : |L(\phi_i) \cap Z(o)| > 0, \ i \in \{1 \ldots n\}\}$$

We can also define the complementary set of the remaining clauses of ϕ *which does not satisfy the previous condition as* $\overline{C}_\phi(o) = \{\phi_i : \phi_i \notin C_\phi(o), \ i \in \{1 \ldots n\}\}$

Definition 6 *Given a preference clause* $\phi_i = l_1 \vee l_2 \vee ... \vee l_m$, *the set* $L(\phi_i) = \{l_1, l_2, ..., l_m\}$ *is the equivalent set-based definition of* ϕ_i *and* $\overline{L}(\phi_i) = \{\neg l_1, \neg l_2, ..., \neg l_m\}$ *is the literal complement set of* $L(\phi_i)$.

Definition 7 *Given an operator* $o \in O$ *of a* STRIPS+P *problem,* $Z(o)$ *is the set of literal defined as:*

$$Z(o) = (Pre(o) \setminus \{p \mid \neg p \in \textit{Eff}(o)^-\}) \cup \textit{Eff}(o)^+ \cup \textit{Eff}(o)^-.$$

Note that the literals in $Z(o)$ hold in any reachable state resulting from the execution of operator o.

To shorten the following definitions and explanation we indicate the state where the operator o is applied with s and the state resulting from the application of o with s'.

Definition 8 *Given an operator* o *and CNF formula* $\phi = \phi_1 \wedge ... \wedge \phi_n$, *we define the set of clauses which will surely be true in the resulting state from the application of* o *as:*

$$C_\phi(o) = \{\phi_i : |L(\phi_i) \cap Z(o)| > 0, \ i \in \{1 \ldots n\}\}$$

We can also define the complementary set of the remaining clauses of ϕ *which does not satisfy the previous condition as* $\overline{C}_\phi(o) = \{\phi_i : \phi_i \notin C_\phi(o), \ i \in \{1 \ldots n\}\}$.

With reference to Definition 8, given a clause $\phi_i = l_1 \vee ... \vee l_{m_i}$ of ϕ, the condition $|L(\phi_i) \cap Z(o)| > 0$ requires that exists at least a literal l_j, with $j \in \{1 \ldots m_i\}$, that, belonging to the set $Z(o)$, will be certainly true in the resulting state from the application of the operator o thus making the clause ϕ true in s'.

Definition 9 *Given an operator* o *and a CNF formula* $\phi = \phi_1 \wedge ... \wedge \phi_n$, *we say that* o ***can make true*** ϕ *if*

1. $|C_\phi(o)| > 0$;

2. for each clause $\phi \notin \overline{C}_\phi(o)$, $\overline{L}(\phi) \not\subseteq Z(o)$.

The first condition in Definition 9 requires that exists at least a clause of the formula which contains some literals which will certainly be true true in the state resulting from the execution of o. The second condition in Definition 9 requires that the remaining clauses of ϕ, which do not belong to $C_\phi(o)$, are certainly not falsified in s', because otherwise the ϕ formula will be certainly falsified in s'.

An operator o that *can make* ϕ does not not guarantees a switch from a state $s \models \neg\phi$ to a state $s' \models \phi$, but it only guarantees $s' \models \phi$ without imposing any condition on the truth value of ϕ in s.

In our compilation scheme of a STRIPS+P problem we have to distinguish, for each kind of preference, different class of operators in order to specialize the operators compilation based on how they interact with the preferences of the problem.

Generally speaking, we can identify three classes of operators regardless of the type of preference considered. An operator is a *threat* for a preference P if in case it is executed it may violate P. An operator is a (potential) *support* for a preference P if in case it is executed it could satisfy P. Finally, an operator o is *neutral* for a preference P if its execution can not influence the current state of the preference. In the following Subsection – we will decline these generic definitions of classes of operators for each type of considered preference providing for each of them a formal definition.

In the following definitions of operators classes we will assume that ϕ (and ψ in the case of dual preferences) are arbitrary CNF formulae $\phi = \phi_1 \wedge ... \wedge \phi_n$ where each clause $\phi_i = l_1 \vee ... \vee l_{m_i}$ for each $i \in \{1...n\}$.

Operators Affecting Always Preferences

An always preference has the following PDDL syntax

$$(preference\ P\ (always\ \phi))$$

where the formula ϕ has to hold in each reached state of the plan. In the following we will abbreviate it with $P = \mathsf{A}_\phi$. According to the semantic provided in Figure **??** a violation of $P = \mathsf{A}_\phi$ is irreversible, when a state s such that $s \not\models \phi$ occurs then it is no longer possible to satisfy P.

Definition 10 *Given an operator o and an always preference $P = \mathsf{A}_\phi$ of a STRIPS+P problem, o is a **violation** of P if there is a clause ϕ_i of ϕ such that:*

1. $\overline{L}(\phi_i) \subseteq Z(o)$;

2. $\overline{L}(\phi_i) \not\subseteq Pre(o)$.

The conditions of Definition 10 require that there exists at least a clause of ϕ (1) whose literals are falsified after the execution of o (2) which is not already false in the state where o is applied.

If an operator violates a preference, the preference is unsatisfied in any state resulting from the application of the operator. The set of always preferences that are violated by an operator o is denoted with $V_\mathcal{A}(o)$.

Definition 11 *Given an operator o and an always preference $P = \mathsf{A}_\phi$ of a STRIPS+P problem, o is a **threat** of P if it is not a violation and there exists a clause ϕ_i of ϕ such that:*

1. $|\overline{L}(\phi_i) \cap Z(o)| > 0$

2. $|L(\phi_i) \cap Z(o)| = 0$

3. $\overline{L}(\phi_i) \not\subseteq Pre(o)$.

The conditions of Definitin 11 require that there exist at least a clause of ϕ that (1) has some (but not all) literals which are falsified after the execution of o, (2) has not literals which are true in the resulting state from the application of o and (3) this clause is not already false in the state where o is applied. The expression $\overline{L}(\phi_i) \not\subseteq Pre(o)$ in Definition 10-11 is necessary to avoids that an operator o is considered a violation/threat when its precondition is already violated in the state where it is applied.

A clause ϕ_i of $P = \mathsf{A}_\phi$ satisfying these conditions is a *threatened clause* of P. A threatened preference (clause) may be falsified by an operator depending on the state where the operator is applied. The set of always preferences threatened by an operator o is denoted with $T_\mathcal{A}(o)$; the set of clauses of an always preference P threatened by o is denoted with $T_\mathcal{A}(o, P)$.

Definition 12 *Given an operator o and an always preference $P = \mathsf{A}_\phi$ of a STRIPS+P problem, o is a **neutral** operator for P if:*

1. for all clauses ϕ_i of P, $|L(\phi_i) \cap Z(o)| > 0$ or $|\overline{L}(\phi_i) \cap Z(o)| = 0$ holds;

2. there exists a clause ϕ_i such that $\overline{L}(\phi_i) \subseteq Pre(o)$.

For example an operator o such that $Pre(o) = \{\neg a\}$ and $Eff(o) = \{\neg b\}$ is a threat for $P1 = \mathsf{A}_{\phi^{P1}}$, a violation for $P2 = \mathsf{A}_{\phi^{P2}}$ and neutral for $P3 = \mathsf{A}_{\phi^{P3}}$ where $\phi^{P1} = c \vee b$, $\phi^{P2} = a \vee b$ and $\phi^{P3} = d$.

Operators Affecting Sometime Preferences

A sometime preference has the following PDDL syntax

$$(preference\ P\ (sometime\ \phi))$$

where the formula ϕ has to become true at least once state in the plan state trajectory. In the following we abbreviate with $P = \mathsf{ST}_\phi$. According to the semantic provided in Figure **??** it is not possible to deliberate if a preference $P = \mathsf{ST}_\phi$ is violated until the end of the planning but the preference can be satisfied at any time during planning if an operator makes ϕ true.

Definition 13 *Given an operator o and a sometime preference $P = \mathsf{ST}_\phi$ of a STRIPS+P problem, o is a **potential support** for P if o can make true ϕ otherwise the operator is **neutral** for P.*

The set of sometime preferences of Π which are potentially supported by the operator o are denoted with $S_{\mathcal{ST}}(o)$.

For example an operator o such that $Pre(o) = \emptyset$ and $Eff(o) = \{\neg b\}$ is a potential support for $P1 = \mathsf{ST}_{\phi^{P1}}$ where $\phi^{P1} = c \vee \neg b$.

Operators Affecting Sometime-before Preferences

A sometime-before preference has the following PDDL syntax and its semantic requires that,

$$(preferences\ P\ (sometime\text{-}before\ \phi\ \psi))$$

whenever ϕ is true in a state s then ψ must have been true in a state before s. In the following we will abbreviate it with $P = \mathsf{SB}_{\phi,\psi}$. Likewise to what exposed for the always preferences a violation of $P = \mathsf{SB}_{\phi,\psi}$ is irreversible, when a state s such that $s \not\models \phi$ occurs without having been made ψ true in a previous state then it is no longer possible to satisfy P.

Speaking informally, if an operator o, applied in a state where ϕ has never been made true before, makes ψ true in the resulting state, then we consider it a *support* for P. Once a support is applied then preference can no longer be violated.

Definition 14 *Given an operator o and a sometime-before preference $P = \mathsf{SB}_{\phi,\psi}$ of a* STRIPS+P *problem, o is a* **potential support** *for P if o can make ψ true.*

An operator o that satisfied Definition 14 is a *potential* support for P because it can act as an *actual* support only under certain preconditions which depend by the stare where it is applied. If a potential support o is applied we can distinguish two possible behaviors:

- if ψ does not become true in the resulting state, then o is *neutral* for P;

- if ψ becomes true in the resulting state and P has not been violated, then o is a *support* for P.

Definition 15 *Given an operator o and a sometime-before preference $P = \mathsf{SB}_{\phi,\psi}$ of a* STRIPS+P *problem, o is a* **threat** *for P if o can make true ϕ.*

Similarly to Definition 14 also in this case the behavior of a threat when it is applied depends by the context. We distinguish the following situations

- if ψ does not become true in the resulting state, then o is *neutral* for P;

- if ψ becomes true in the resulting state and the formula ϕ has become true at least once in a earlier state, than s is *neutral* for P otherwise if the formula ϕ has never become true in previous states, then o is a *violation* for P.

The set of sometime-before preferences of Π which are threatened and potentially supported by the operator o are denoted respectively with $T_{\mathcal{SB}}(o)$ and $S_{\mathcal{SB}}(o)$.

Definition 16 *Given an operator o and a sometime-before preference $P = \langle \phi, \psi \rangle$ of a* STRIPS+P *problem, o is a* **neutral** *for P if o is not a support or a threat.*

For example an operator o such that $Pre(o) = \emptyset$ and $Eff(o) = \{b\}$ is a potential support for $P1 = \mathsf{SB}_{\phi^{P1},\psi^{P1}}$, a threat for $P2 = \mathsf{SB}_{\phi^{P2},\psi^{P2}}$ and neutral for $P3 = \mathsf{SB}_{\phi^{P3},\psi^{P3}}$ where $\phi^{P1} = c$ and $\psi^{P1} = a \vee b$ $\phi^{P2} = c \vee b$ and $\psi^{P2} = d$ and $\phi^{P3} = d$ and $\psi^{P3} = e$.

Operators Affecting At-most-once Preferences

An at-most-once preference has the following PDDL syntax

$$(preference\ P\ (at\text{-}most\text{-}once\ \phi)$$

where the formula ϕ has to become true at most once in the plan state trajectory. In the following we will abbreviate it with $P = \mathsf{AO}_\phi$.

Definition 17 *Given an operator o and an at-most-once preference $P = \mathsf{AO}_\phi$ of a* STRIPS+P *problem, o is a* **threat** *operator for P if o can make true ϕ.*

We distinguish the following situations:

- if ϕ has never become true in all the states earlier than the state s where o is applied and becomes true in the resulting state from the application of o then the operator behaves as *neutral* for P;

- if ϕ does not become true in the state resulting state from the application of o, then o behaves as *neutral* for P;

- if ϕ has become true in a earlier state and will become true in the resulting state from the application of o, the o behaves as a *violation* for P.

The set of at-most-once preferences of Π which are threatened by the operator o are denoted with $T_{\mathcal{AO}}(o)$.

For example an operator o such that $Pre(o) = \emptyset$ and $Eff(o) = \{\neg b\}$ is a potential support for $P1 = \mathsf{AO}_{\phi^{P1}}$ where $\phi^{P1} = c \vee \neg b$.

Compilation of Qualitative Preferences

Before describing the general compilation scheme, some further definitions should be provided.

Definition 18 *Given a* STRIPS+P *problem Π if an operator o of Π is neutral for every given preference of Π over $\mathcal{A}$, $\mathcal{SB}$, $\mathcal{ST}$, $\mathcal{AO}$ and $\mathcal{SG}$ then we say that o is* **neutral** *for Π. The set of all the neutral operators for Π is denoted by $O_{neutral}$.*

Definition 19 *Given an operator o of a* STRIPS+P Π *problem the set $P_{affected(o)}$ of prefereces affected by o is defined as:*

$$P_{affected(o)} = T_{\mathcal{A}}(o) \cup T_{\mathcal{SB}}(o) \cup S_{\mathcal{SB}}(o) \cup T_{\mathcal{AO}}(o) \cup S_{\mathcal{ST}}(o).$$

Given a STRIPS+P problem, an equivalent STRIPS+ problem can be derived by translation which has some similarities to what proposed by Keyder and Geffner for soft goals but also significant difference. The scheme proposed by Keyder and Geffner is considerable simpler than ours because it does not to consider the interaction between actions and preferences such as threats, supports and violations. In order to simplify the compilation scheme we don't consider the compilation of soft goals because it can be easily added using the same method of Keyder and Geffner.

Moreover we assume that every always and sometime-before preference are not violated and every sometime preference are not satisfied in the problem initial state I. Before starting the compilation, we carry out the following checks in order to exclude some preferences from the process:

- for each $P = \mathsf{A}_\phi \in \mathscr{P}_\mathcal{A}$ we check that $I \models \phi$, if the condition does not hold we exclude P from the compilation increasing the cost of the plan by $u(P)$ because it is already violated in I;

- for each $P = \mathsf{SB}_{\phi,\psi} \in \mathscr{P}_{\mathcal{SB}}$ we check that $I \not\models \phi$, if the condition does not hold we exclude P from the compilation increasing the cost of the plan by $u(P)$ because it is already violated in I; after that we check that $I \not\models \phi \wedge I \models \psi$, if the condition hold we exclude P because it is already satisfied in I;

- for each $P = \mathsf{ST}_\phi \in \mathscr{P}_{\mathcal{ST}}$ we check that $I \models \phi$, if the condition holds we exclude P because it is already satisfied in I;

Given a STRIPS+P problem $\Pi = \langle F, I, O, G, \mathscr{P}, c, u \rangle$, the compiled STRIPS+ problem of Π is $\Pi' = \langle F,' I', O', G', c' \rangle$ where:

- $F' = F \cup V \cup D \cup C \cup \overline{C'} \cup \{\textit{normal-mode}, \textit{end-mode}\}$;

- $I' = I \cup \overline{C'} \cup V_{\mathcal{ST}} \cup S_{\mathcal{AO}} \cup \{\textit{normal-mode}\}$;

- $G' = G \cup C'$;

- $O' = \{\textit{collect}(P_i), \textit{forgo}(P_i) \mid P_i \in \mathscr{P}\} \cup \{\textit{end}\} \cup \{\textit{comp}(o, \mathscr{P}) \mid o \in O\}$

- $c'(o') = \begin{cases} u(P) & \text{if } o' = \textit{forgo}(P) \\ c(o') & \text{if } o' = \textit{comp}(o, \mathscr{P}) \\ 0 & \text{otherwise} \end{cases}$

where:

- $V = \cup_{i=1}^{|\mathscr{P}|} \{P_i\text{-}\textit{violated}\}$;

- $V_{\mathcal{ST}} = \cup_{i=1}^{|\mathscr{P}_{\mathcal{ST}}|} \{P_i\text{-}\textit{violated}\}$, $\mathscr{P}_{\mathcal{ST}} \subseteq \mathscr{P}$; this is the set of the violation predicates restricted to the subset of the sometime preferences of all preferences $\mathscr{P}$;

- $S_{\mathcal{AO}} = \cup_{i=1}^{|\mathscr{P}_{\mathcal{AO}}(I)|} \{P_i\text{-}\textit{seen}\}$, $\mathscr{P}_{\mathcal{AO}}(I) = \{P_i = \mathsf{AO}_\phi \mid P_i \in \mathscr{P}_{\mathcal{AO}}, I \models \phi\}$; this is the set of *seen* predicates (see Section) for those at-most-once preferences such that $I \models \phi$;

- $C' = \{P' \mid P \in \mathscr{P}\}$ and $\overline{C'} = \{\overline{P'} \mid P \in \mathscr{P}\}$;

- $\textit{forgo}(P_i) = \langle \{\textit{end-mode}, P_i\text{-}\textit{violated}, \overline{P_i'}\}, \{P', \neg \overline{P'}\} \rangle$;

- $\textit{end} = \langle \{\textit{normal-mode}\}, \{\textit{end-mode}, \textit{normal-mode}\} \rangle$;

- $\textit{comp}(o, \mathscr{P})$, which definition is provided in Definition 20;

Forgo and Collect Actions For each preference P the transformation of Π into Π' adds a dummy hard goal P' to Π' which can be achieved by two ways: with action $\textit{collect}(P)$, that has cost 0 but requires P to be satisfied (i.e. P-*violated* is false in the goal state for all kinds of preferences except for sometime), or with action $\textit{forgo}(P)$, that has cost equal to the utility of P and can be performed only if P is false (P-*violated* is true in the goal state). For each preference, exactly one of $\textit{collect}(P)$ and $\textit{forgo}(P)$ appears in the plan.

Operator Compilation Function The function $\textit{comp}(o, \mathscr{P})$ which transforms an original operator o into the equivalent compiled one is splitted in two parts. If the operator is neutral ($o \in O_{neutral}$) then the function just extends $Pre(o)$ with the predicate *normal-mode* in order to scorporate the execution of the domain operators and the evaluationof *forgo* and *collect* actions at the end of the planning.

If the operator o is not neutral ($o \in O - O_{neutral}$) where $|P_{affected}(o)| = n$, then the compilation function $\textit{comp}(o, \mathscr{P})$ extends its effects, for each affected preference P_i, by adding a set of conditional effects denoted with $\mathcal{W}(o, P_i)$ whose definition depends by o, the class of preference P_i belongs to and the way how o interacts with P_i. In Subsections – we will detail the definition of $\mathcal{W}(o, P_i)$ for each class of preference.

But we want a compiling problem which belongs to STRIPS+ class and so we have to compile away the conditional effects (see Section).

Extension of the initial state Note that the original initial state I is extended to I' with the set of literals $V_{\mathcal{ST}}$, which contains the literal P_i-*violated* for each sometime preference of the problem. The literal P_i-*violated* states that the related preference P_i is (temporarily) violated in I until a support operator for P is applied.

Furthermore the original initial state I is extended with the set of literals $S_{\mathcal{AO}}$ which contains, for each at-most-once preference $P_i = \mathsf{AO}_{\phi_i}$ of the problem such that $I \models \phi_i$, the literal *seen-P_i*. This additional fluent *seen-P_i* states that the formula ϕ_i, involved in P_i, is satisfied in I. This precaution is necessary to correctly capture any possible violations of at-most-once preferences.

Definition 20 *Given an operator o the corresponding compiled operator is defined using the following function:*

$$Pre(o') = Pre(o) \cup \{\textit{normal-mode}\}$$
$$Eff(o') = Eff(o) \cup \bigcup_{P_i \in P_{affected}(o)} \mathcal{W}(o, P_i)$$

where $\mathcal{W}(o, P_i)$ is the set of conditional effects concerning the affected preference P_i. If $o \in O_{neutral}$ then $Eff(o') = Eff(o)$.

Compilation of an Always Preference

We now present the transformation of operators that threaten or violate a preference in class $\mathcal{A}$.

Before defining the extending effects used to compile an operator o which affects an always preference, we introduce some useful notation in order to simplify the formalisation. For an operator o and a preference clause ϕ_i:

- $NA(o)_{\phi_i} = \{l_j \in L(\phi_i) \mid \neg l_j \in (Eff(o)^+ \cup Eff(o)^-)\}$ is the set of literals in $L(\phi_i)$ falsified by the effects of o;

- $AA(o)_{\phi_i} = L(\phi_i) \setminus NA(o)_{\phi_i}$ is the set of literals in $L(\phi_i)$ *not* falsified by the effects of o;

- $\overline{AA}(o)_{\phi_i}$ is the literal-complement set of $AA(o)_{\phi_i}$.

Definition 21 *Given an always preference* $P = \mathsf{A}_\phi$ *and an operator* o *which affects* P*, the conditional effect set* $\mathcal{W}(o, P)$ *in the compiled version* o' *of* o *(according to Definition 20) is defined as:*

$$\mathcal{W}(o, P) = \begin{cases} \{when\ (cond(o, P))\ (P\text{-}violated)\} \\ \quad \textit{if } o \textit{ is a threat for } P \\ \{when\ (\text{TRUE})\ (P\text{-}violated)\} \\ \quad \textit{if } o \textit{ is a violation for } P \end{cases}$$

where:

- $cond(o, P) = \bigvee_{\phi_i \in T_\mathcal{A}(o,P)}(l_1 \wedge \ldots \wedge l_q)$, $\{l_1, \ldots, l_q\} = \overline{AA}(o)_{\phi_i}$.

For each preference $P \in P_\mathcal{A}$ affected by an operator o, the compiled operator o' contains a conditional effect whose effect $P\text{-}violated$ and conditions depends on how o affects P. If o is a violation for P, then the condition is always true (i.e. it is the special literal TRUE that holds in every state). Instead, if o is a *threat* to P, the effects are extended with a conditional described below.

If o is a threat for P the condition $(cond(o, P))$ requires $\overline{AA}(o)_{\phi_i}$ to hold for at least one $\phi_i \in T_\mathcal{A}(o, P)$, which implies that ϕ is false in the state generated by o'.

Compilation of a Sometime Preference

We now present the transformation of an operator that is a potential supports for a preference P in class ST.

Definition 22 *Given a sometime preference* $P = \mathsf{ST}_\phi$ *and an operator* o *which potentially supports* P*, the conditional effect set* $\mathcal{W}(o, P)$ *in the compiled version* o' *of* o *(according to Definition 20) is defined as:*

$$\mathcal{W}(o, P) = \{when(cond_\mathcal{S}(o, P))\ (\neg P\text{-}violated)\}$$

where:

- $cond_\mathcal{S}(o, P) = \{\phi_i \mid \phi_i \in \overline{C}_\phi(o)\}$.

As specified above, the general compilation scheme introduces for each problem preference $P = \mathsf{ST}_\phi$ $P \in \mathscr{P}_{\mathcal{ST}}$ a predicate $P\text{-}violated$ in the compiled initial state I' if ϕ does not hold in the original problem initial state (otherwise the preference is considered satisfied and therefore not compiled).

This is necessary because, according to the semantics of the preference in class ST, a plan π' in the compiled problem satisfies ST_ϕ iif ϕ is true at least once in the state trajectory of π'.

These $P\text{-}violated$ predicates are falsified by the operators that could make ϕ true in the state s' resulting from this application. Such operators are potential support that we have defined in Section .

An operator o could make a formula ϕ true when there are some clauses of ϕ that will surely be true in the state s' resulting from the application of o. So, if all $\phi_i \in \overline{C}_\phi(o)$ hold in s (where $\overline{C}_\phi(o)$ is the set of clauses of ϕ that are not surely true in s' - see Definition 8), i.e. $cond(o, P)$ holds in

s, then ϕ will be true in s', and $P\text{-}violated$ should be falsified by the effects of o'.

In order to satisfy a given preference ST_ϕ, which is temporarily violated in I', a potential support o for P has to be inserted in the plan and applied in a state where $(cond_\mathcal{S}(o, P))$ holds, making an actual support for P. Otherwise, operator o behaves as a neutral operator for P, leaving the preference violated in s'.

Compilation of a Sometime-before Preference

We now present the transformation of operators that potentially support or threat a preference in $\mathcal{SB}$.

Definition 23 *Given a sometime-before preference* $P = \mathsf{SB}_{\phi,\psi}$ *and an operator* o *which affects* P*, the conditional effect set* $\mathcal{W}(o, P)$ *in the compiled version* o' *of* o *(according to Definition 20) is defined as:*

$$\mathcal{W}(o, P) = \begin{cases} \{when\ (cond_\mathcal{S}(o, P))\ (seen\text{-}\psi)\} \\ \quad \textit{if } o \textit{ is a potential support for } P \\ \{when\ (cond_\mathcal{T}(o, P))\ (P\text{-}violated)\} \\ \quad \textit{if } o \textit{ is a threat for } P \\ \{when\ (cond_\mathcal{S}(o, P))\ (seen\text{-}\psi), \\ \quad when\ (cond_\mathcal{S}(o, P))\ (P\text{-}violated)\} \\ \quad \textit{if } o \textit{ is a both a threat and support for } P \end{cases}$$

where:

- $cond_\mathcal{S}(o, P) = \{\psi_i \mid \psi_i \in \overline{C}_\psi(o)\}$
- $cond_\mathcal{T}(o, P) = \{\neg seen\text{-}\psi\} \cup \{\phi_i \mid \phi_i \in \overline{C}_\phi(o)\}$.

The definition of the effects used to extend the effect of an operator o which affects a sometime-before preference P in Definition 23 depends on the class of operators which o belongs. We have to distinguish if o is a potential support, a threat or both for P. Remember that the semantics of the preferences belonging to $\mathcal{SB}$ requires that a preference $\mathsf{SB}_{\phi,\psi}$ is satisfied by a state-trajectory if, whenever ψ becomes true in a state then ϕ must have become true in a previous state.

If o is a potential support for P then this operator can behave in two different ways when it is executed. Recalling Definition 9, an operator o could make true a formula ϕ when there exists some clauses of ϕ that will surely be true in the state resulting from the application of o. So, if all $\psi_i \in \overline{C}_\psi(o)$ hold in s (where $\overline{C}_\psi(o)$ is the set of clauses of ψ that are not surely true in s' - see Definition 8) , i.e. $cond_\mathcal{S}(o, P)$ holds in s, then ψ will be true in s' and then we have to keep track of this fact by making the predicate $seen\text{-}\psi$ true in the effects $\textit{Eff}_\mathcal{T}(o, P)$ of o'.

If o is a threat for P its compilation is similar. A threat for $\mathsf{SB}_{\phi,\psi}$ behaves as violation in the case that the operator makes ϕ true in s' (i.e., when all $\phi_i \in \overline{C}_\phi(o)$ holds in s, which is the condition of $cond_{\mathcal{T}(o,P)}$), and that the ψ has never been made true in the states preceeding s. If both these conditions, specified in $cond_\mathcal{T}(o, P)$, hold in s then predicate $P-violated$ is included in the effects of o'.

We have also to consider the case in which an operator o is both a threat and a support for P. In this case o can

behaves in the following ways: making ϕ true in s', making ψ true in s' and making both ϕ and ψ true in s'. In order to handle these situations, the compiled operator o' contains both conditional effects of o as threat of P and as support of P. Note that this correctly captures the violation of P determined by ϕ and ψ becoming simultaneously true by execution of o.

Compilation of an At-Most-Once Preference

We now present the transformation of an operator that threats a preference in $\mathcal{AO}$.

Definition 24 *Given an at-most-once preference $P = \mathsf{AO}_\phi$ and an operator o which affects P, the conditional effect set $\mathcal{W}(o, P)$ in the compiled version o' of o (according to Definition 20) is defined as:*

$$\mathcal{W}(o,P) = \{$$
$$\quad when \ (cond_N(o, P)) \ (seen\text{-}\phi))$$
$$\quad when \ (cond_T(o, P))) \ (P\text{-}violated)\}$$

where:

- $cond_N(o, P) = \{\neg seen\text{-}\phi\} \cup \{\phi_i \mid \phi_i \in \overline{C}_\phi(o)\}$
- $cond_T(o, P) = \{seen\text{-}\phi\} \cup \{\phi_i \mid \phi_i \in \overline{C}_\phi(o)\} \cup$
$$\{\textstyle\bigvee_{\phi_i \in C_\phi(o)} (\neg l_1 \wedge ... \wedge \neg l_q) \mid$$
$\{l_1, ..., l_q\} = L(\phi_i)\}.$

The semantic of a preference in class $\mathcal{AO}$ (Figure **??**) requires that a preference $P = \mathsf{AO}_\phi$ is satisfied by the state-trajectory generated by a plan π if ϕ *becomes* true in a state s' at most once during the execution of the plan. A formula ϕ becomes true in a state s' due to the execution of an operator o applied in a state s iif $s \models \neg\phi$ and $s' \models \phi$.

If ϕ holds in the problem initial state I, then it is required that either it stay true until the end of the plan. or it becomes false at some successor state of I in the state-trajectory generated by the plan and it then stays always false.

If an operator o can make ϕ true for the first time in the plan trajectory of a plan, then it behaves as a neutral operator for P. On the other hand, if o can make ϕ true after having been true and become false in past states of the trajectory, then o behaves as a threat for P.

In order to correctly capture the possible violation of P, the set of effects extending o $\mathcal{W}(o, P)$ has two conditional effects. The first one is a neutral effect that is used to catch the behavior of o as neutral operator for P. Condition $cond_N(o, P)$ requires that ϕ has been never changed truth value from true to false in preceeding state using the negated predicate $\neg seen\text{-}\phi$ and that ϕ will be true in s' using the condition $\{\phi_i \mid \phi_i \in \overline{C}_\phi(o)\}$ similarly to what done in the compilation schemes previously presented. If the conditions specified in $cond_N(o, P)$ hold in the state s where o is applied, then we take into account that ϕ becomes true in s' for the first time stating in the effects of the neutral conditional effect predicate $seen\text{-}\phi$.

The second conditional effect is a violating effect that is used to catch the the behavior of of o as a threat for P.

This happens when the following conditions, specified in $cond_{T(o,P)}$ hold: (specified in $cond_{\mathcal{V}(o,P)}$):

- ϕ has alreay been made true in a state preceding s; this is expressed by the predicate $seen\text{-}\phi$ when o is applied;

- ϕ is made true in the resulting state s'; this is guaranteed by the conditions $\{\phi_i \mid \phi_i \in \overline{C}_\phi(o)\}$;

- ϕ is false in the state s where o is applied; this is specified by requiring that at least a clause in $C_\phi(o)$ is false in s.

If all these conditions hold in the state where o is applied then P will be violated in the state resulting from the application of o.

Compilation of Conditional Effects

As described before, given an operator o of a STRIPS+P problem which affects a set of n preferences, the corresponding compiled operator should have in its effects a set $\{(when \ c_i \ e_i) \mid i = 1 \ldots m\}$ of $m \leq 2n$ conditional effects, which are built using the compilation schema described in the previous subsections.

In order to keep the compiled problem in the STRIPS+ class, the conditional effects of o' should be compiled away by replacing o' with an equivalent set of operators without conditional effects. In the literature, there are two main general methods for generating this equivalent set of unconditional operators.

The first method, introduced by Gazen and Knoblock [**?**], works by recursively splitting o' for each conditional effect *(when c_i e_i)* into a couple of new operator, o'' and $\overline{o}''$, such that:

$$pre(o'') = pre(o') \cup \{c_i\}$$
$$eff(o'') = eff(o') \cup e_i$$
$$pre(\overline{o}'') = pre(o') \cup \{\neg c_i\}$$
$$eff(\overline{o}'') = eff(o').$$

This method is not practicable because it leads to an exponential blow up of operators (in our case $O(2^m)$ for each operator affecting n preferences), but the compiled plan preserves exactly the length of the original plan.

The second method, proposed by Nebel [**?** , see proof of Theorem 20], generates a polynomial number of new operators, but it increases polynomially the plan length. The main idea is to simulate the parallel behaviour of the conditional effects of an operator by a replacing it with an equivalent sequence of unconditional operators. For each operator o' with m conditional effects, Nebel's schema introduces m pairs of new operators, that separately evaluate the condition c_i of each conditional effect *(when c_i e_i)* and possibly "activates" the corresponding effect e_i. One of the operators in the pair for *(when c_i e_i)* contains precondition c_i and an effect indicating that e_i is activated, while the other, that is mutex with the first, contains precondition $\neg c_i$ and does not activate e_i. In order to avoid possible (positive or negative) interference in the sequentialisation of the conditional effects through the new operators (e.g., if *(when c_1 e_1)* and *(when c_2 e_2)* are conditional effects of o' and $e_1 \models c_2$), the activated effects in the operator sequence are not made immediately true, but

they are deferred to the end of the sequence (i.e., after all conditional effects conditions have been evaluated). This is done by using an additional set of operators, called "copying operators" in [**?**], which copy the activated effects to the state description after all operators in the sequence have been executed (For more details the reader is referred to [**?**].)

In order to deal with the conditional effects generated by our compilation of PDDL3 preferences, we have implemented and used Nebel's compilation method because a compiled operator can in principle contain many conditional effects, which makes Gazen and Knoblock's method impractical given its exponential complexity. Moreover, the conditional effects needed to compile PDDL3 preference have a particular structure that allows us to simplify and optimize Nebel's general method. In particular we would like avoid the so-called "copying operators" operators maintaining the semantics of conditional effects (which requires that all conditions are evaluated at the same time). First of all, note that the conditional effects that refer to different affected preferences can not interfere with each other because they involve different fluents.

After that we have to pay attention to those class of preferences whose compilation introduces more than one conditional effect, because, affecting the same fluents, they can generate interference. In our previous discussion there are two classes of preferences having this feature, i.e. sometime-before and at-most-once preferences (see Subsections and). In the first case, the set of problematic conditional effects refers to those operators which threats and supports a sometime-before preference at the same time, in the second case to those operators that threats at-most-once preference.

Concerning at-most-once preferences, an interference could arise through the predicate $seen\text{-}\phi$, that is both an effect of the first conditional effect and a condition of the second. However, the conditional effect interference disappears, if in the compilation, the pair of unconditional operators for $(when\ (cond_\mathcal{V}(o, P))\ (P-violated))$ is constrained to be ordered before the other pair. If we evaluated these conditional effects without following this order, the execution would be equivalent to not evaluating all conditional effects of the same operator simultaneously, thus risking to recognize a violation even if this does not happen. Indeed, given a preference $P = \mathsf{AO}_\phi$ and a threat operator o, if ϕ becomes true for the first time in s' after the application of o, an we check the condition $cond_\mathcal{N}(o, P)$ before $cond_\mathcal{V}(o, P)$, then we could detect a violation that may not have happened.

We can expose similar considerations for sometime-before preferences. In this case, if we do not check the condition $cond_\mathcal{V}(o, P)$ before $cond_\mathcal{T}(o, P)$, we risk to not correctly identifying a possible violation if ϕ and ψ become true at the same time.

Consequently, starting from the previous observations which show that it is possible to eliminate any interferences within our context, Nebel's copying operators are not needed in the compilation of our conditional effects; furthermore, in the compiled problem, we can force an arbitrary total order of the unconditional operator pairs, paying attention that the ordering constraints dealing with the potential interference

between the conditional effects arising from at-most-once and sometime-before preferences are satisfied.

These changes to Nebel's schema simplify it, and have some beneficial consequences: (1) the compiled problem has smaller size in terms of number of operators, and (2) the search effort of a planner can be reduced because the solution plans are shorter without the coping actions,[3] and because the sequence of the unconditional operators can be explicit in the compiled problem, while with the original compilation it is built at search time by the planner.

Since our compilation of conditional effects can be easily derived from Nebel's method, instead of giving all the formal details of the translation, we illustrate the compilation with an example, referring the reader to [**?**] for a formal description. Moreover, in [**?**] we give a detailed description of the final (without conditional effects) compiled problem for any STRIPS+ problem with always preferences.

Compilation Example

Consider the following operator $o = \langle Pre(o), \mathit{Eff}(o)\rangle = \langle \varnothing, \{a, \neg c\}\rangle$ with cost κ, which affects two preferences, $P1 \in \mathcal{A}$ and $P2 \in \mathcal{AO}$, where $P1 = (always\ \phi^{P1})$ and $P2 = (at\text{-}most\text{-}once\ \phi^{P2})$ and $\phi^{P1} = \phi_1^{P1} \wedge \phi_2^{P1} = (a \vee b) \wedge (c \vee d)$ and $\phi_1^{P2} \wedge \phi_2^{P2} = (\neg c \vee e) \wedge (d \vee f)$.

Compilation of P1   Using the conditions specified in Definition 11, we can say that o is a *threat* for $P1$ because there exists a clause $\phi_2^{P1} = (c \vee d)$ of ϕ^{P1} such that:

1. it contains at least a literal, which is c, that is negated by the effects of o: $|\overline{L}(\phi_1^{P2}) \cap Z(o)| = |\overline{L}(c \vee d) \cap Z(o)| = |\{\neg c, \neg d\} \cap \{a, \neg c\}| = |\{\neg c\}| > 0$;

2. the other literal in the clause, which is d, is not made true in the resulting state from the application of o: $|L(\phi_1^{P2}) \cap Z(o)| = |L(c \vee d) \cap Z(o)| = |\{c, d\} \cap \{a, \neg c\}| = 0$;

3. (the clause) is not already negated in the state where o is applied because its negation is not implied by the precondition of o: $\overline{L}(c \vee d) = \{\neg c, \neg d\} \nsubseteq Pre(o)$.

Operator o threatens only one clause of $P1$, i.e. $T_\mathcal{A}(o, P1) = \{\phi_2^{P1}\} = \{c \vee d\}$ (remember that $T_\mathcal{A}(o, P1)$ is the set of threatened clauses of $P1$ by o), which could be falsified if o is applied.

Using the preliminary definitions provided in Section , we define the set $\overline{A\!A}(o)_{\phi_2^{P1}} = \{\neg d\}$ which contain those literal of the threatened clause ϕ_2^{P1} that are not falsified by o.

In this case o, denying the literal c, threatens a single clause of $P1$, i.e. $\phi_2^{P1} = c \vee d$, and therefore we have to check the literal d is true in the precondition to capture the possible violation.

Starting from these considerations and using Definition 21 we define the conditional effect $\mathcal{W}(o, P1)$ to add to

---

[3]Given a plan π for a problem with conditional effects and the corresponding plan π' for the compiled problem, we have $|\pi'| \leq |\pi| * m$ while with Nebel's general method this bound is $|\pi'| \leq |\pi| * (3 + m)$ [**?**].

Eff(o):

$$\mathcal{W}(o, P1) = \{when \ (\neg d) \ (P1\text{-}violated)\}.$$

Compilation of P2 According to Definition 17, o is also a *threat* for $P2$ because it could make true ϕ^{P2} in the resulting state s' from its application; indeed there exists a clause of ϕ^{P2}, i.e. $\phi_1^{P2} = c \vee e$, such that:

- it contains at least a literal, i.e. $\neg c$, that will be surely true in s': $|L(\phi_1^{P2}) \cap Z(o)| = |L(\neg c \vee e) \cap Z(o)| = |\{\neg c, e\} \cap \{a, \neg c\}| = |\{\neg c\}| > 0$

- the other clause, which is $\phi_2^{P2} = d \vee f$, is not negated by the execution of o: $|\overline{L}(\phi_2^{P2}) \cap Z(o)| = |\overline{L}(d \vee f) \cap Z(o)| = |\{\neg d, \neg f\} \cap \{a, \neg c\}| = 0$

We denote with $C_{\phi^{P2}}(o) = \{\phi_1^{P2}\} = \{\neg c \vee e\}$ the set of clauses of ϕ^{P2} that will be surely true in the resulting state s' and with $\overline{C}_{\phi^{P2}}(o) = \{\phi_2^{P2}\} = \{d \vee f\}$ the set of clauses of ψ the remaining ones. With reference to Section , we define the set of conditional effects related to the threatened at-most-once preference $P2$ as:

$$\mathcal{W}(o, P_2) = \{$$
$$when \ (cond_N(o, P_2)) \ (seen\text{-}\phi^{P2}))$$
$$when \ (cond_T(o, P_2))) \ (P\text{-}violated)\}$$

where:

- $cond_N(o, P2) = \{\neg seen\text{-}\phi^{P2}\} \cup \{\phi_i \mid \phi_i \in \overline{C}_{\phi^{P2}}(o)\} = \{\neg seen\text{-}\phi^{P2}\} \cup \{\phi_2^{P2}\} = \{\neg seen\text{-}\phi^{P2}, d \vee f\}$

- $cond_T(o, P2) = \{seen\text{-}\phi^{P2}\} \cup \{\phi_i \mid \phi_i \in \overline{C}_{\phi^{P2}}(o)\} \cup \{\bigvee_{\phi_i \in C_{\phi^{P2}}(o)}(\neg l_1 \wedge ... \wedge \neg l_q) \mid \{l_1, ..., l_q\} = L(\phi_i)\} = \{seen\text{-}\phi^{P2}\} \cup \{\phi_2^{P2}\} \cup \{\{\neg \neg c \wedge \neg e\} \mid \{\neg c, e\} = L(\phi_1^{P2})\} = \{seen\text{-}\phi^{P2}, d \vee f, c \wedge \neg e\}$

The condition $cond_N(o, P2)$ to activate the fluent $seen\text{-}\phi^{P2}$ requires that the following condition hold in the state where o is applied:

- the fluent $seen\text{-}\phi^{P2}$ has to be false and so ϕ^{P2} never became true until then;

- those clauses of $P2$ which are not affected by o, i.e. $\phi_2^{P2} = d \vee f$, have to be true making sure that ϕ^{P2} becomes true.

The condition $cond_T(o, P_2))$ to activate the violation of $P2$ requires that:

- the fluent $seen\text{-}\phi^{P2}$ has to be true and so ϕ^{P2} has already been made true;

- those clauses of $P2$ which are not affected by o, i.e. $\phi_2^{P2} = d \vee f$, have to be true and those clauses of $P2$ which are affected by o, i.e. $\phi_1^{P2} = \neg c \vee e$, have to be false ensuring that ϕ^{P2} passes from false to true.

Compilation of conditional effects According to what described above we can say that o it is simultaneously a threat to both $P1$ and $P2$ and we denote the set of affected preference by o as:

$$P_{affected}(o) = \{P1, P2\}.$$

The compiled operator o' with conditional effects is defined, according to the compilation scheme provided in Section **??** and using Definition 20, as:

$Pre(o') = Pre(o) \cup \{normal\text{-}mode\}$
$Eff(o') = Eff(o) \cup \mathcal{W}(o, P1) \cup \mathcal{W}(o, P2) = Eff(o) \cup \{when \ (\neg b) \ (P1\text{-}violated)\} \cup \{when \ (\neg seen\text{-}\phi^{P2} \wedge (d \vee f)) \ (seen\text{-}\phi^{P2}),$
$when \ (seen\text{-}\phi^{P2} \wedge (d \vee f) \wedge (c \wedge \neg e)) \ (P2\text{-}violated)\}.$

We compile away the three conditional effects replacing o' with the set of operators defined as:

$$O_{P_{affected}(o)} = \{o_{P1}, \overline{o}_{P1}, o_{P2}^V, \overline{o}_{P2}^V, o_{P2}^N, \overline{o}_{P2}^N\},$$

which contains, for each conditional effect specified in $Eff(o')$, a pair of mutually exclusive unconditional operators. In particular:

- o_{P1} and $\overline{o}_{P1}$ refer to the first conditional effect of o' introduced by the compilation to handle the possible violation of P_1;

- o_{P2}^V and $\overline{o}_{P2}^V$ refer to the second conditional effect of o' introduced by the compilation to handle the possible violation of P_2;

- o_{P2}^N, $\overline{o}_{P2}^N$ refer to the third conditional effect of o' introduced in order to correctly updating the predicate $seen\text{-}\phi^{P2}$ when ϕ^{P2} becomes true.

These six operators are defined in order to guaratee that, in any state s where o can be applied affecting $P1$ and $P2$, the sequence $\omega_{P_{affected}(o)}$ of three operators in $O_{P_{affected}(o)}$, defined as follows, can be applied:

$$\omega_{P_{affected}(o)} = \langle o'_{P1}, o'^V_{P2}, o'^N_{P2} \rangle$$

where:

- $o'_{P1} = \overline{o}_{P1}$ if o violates $P1$ when applied in s, $o'_{P1} = o_{P1}$ if o does not violate $P1$ when applied in s;

- $o'^V_{P2} = \overline{o}_{P2}^V$ if o violates $P2$ when applied in s, $o'^V_{P2} = o_{P2}^V$ if o does not violate $P2$ when applied in s;

- $o'^N_{P2} = \overline{o}_{P2}^N$ if o makes ψ true when applied in s, $o'^N_{P2} = o_{P2}^N$ if o does not make ψ true when applied in s.

These operators are defined in the following way:

- operators o_{P1} and o_{P2}:

$Pre(o_{P1}) = Pre(o) \cup \{\neg pause, b\}$
$Eff(o_{P1}) = \{P1\text{-}done, pause\}$

$Pre(\overline{o}_{P1}) = Pre(o) \cup \{\neg pause, \neg b\}$
$Eff(\overline{o}_{P1}) = \{P1\text{-}done, P1\text{-}violated, pause\}$

- operators o_{P2}^V and $\overline{o}_{P2}^V$:

  $Pre(o_{P2}^V) = \{P1\text{-}done, pause\} \cup \{\neg((seen\text{-}\phi^{P2}) \wedge (\neg c \vee e) \wedge (d \vee f))\}$
  $Eff(o_{P2}^V) = \{P2\text{-}done, \neg P1\text{-}done\}$

  $Pre(\overline{o}_{P2}^V) = \{P1\text{-}done, pause, seen\text{-}\phi^{P2}, \neg c \vee e, d \vee f\}$
  $Eff(\overline{o}_{P2}^V) = \{P2\text{-}done, \neg P1\text{-}done, P2\text{-}violated\}$

- operators o_{P2}^N and $\overline{o}_{P2}^N$:

  $Pre(o_{P2}^N) = \{P2\text{-}done, pause\} \cup \{\neg(\neg seen\text{-}\phi^{P2} \wedge (d \vee f))\}$
  $Eff(o_{P2}^N) = \{\neg P2\text{-}done, \neg pause\}$

  $Pre(\overline{o}_{P2}^N) = \{P2\text{-}done, pause, \neg seen\text{-}\phi^{P2}, d \vee f\}$
  $Eff(\overline{o}_{P2}^N) = \{\neg P2\text{-}done, \neg pause, seen\text{-}\phi^{P2}\}$

The operators of $O_{P_{\textit{affected}}(o)}$ involve additional predicates that are used for two purposes: (1) evaluating the conditional effects of o' in sequence and (2) avoiding that the domain state changes because of the execution of domain operator during the evaluation.

The $P_i\text{-}done$ predicates force the planner to strictly follow a fixed order avoiding repetitions in $\omega_{P_{\textit{affected}}(o)}$. Once the planner starts the sequence $\omega_{P_{\textit{affected}}(o)}$, no other operator, except those in $O_{P_{\textit{affected}}(o)}$, can be applied before the application of the sequence is completed.

The order of execution may be arbitrary, except for those operators concerning the violation of $P2$, i.e. o_{P2}^V and $\overline{o}_{P2}^V$, whose evaluation takes place before that of the operators o_{P2}^N and $\overline{o}_{P2}^N$, appointed for the update the $seen\text{-}\phi^{P2}$ predicate. As described previously, this expedient is necessary in order to correctly capture the violations of at-most-once preference when they occour.

Concerning the cost of the compiled operators of $O_{P_{\textit{affected}}(o)}$ we have $c(o_{P1}) = c(\overline{o}_{P1}) = cost(o) = \kappa$ and $c(o_{P2}^V) = c(\overline{o}_{P2}^V) = c(o_{P2}^N) = c(\overline{o}_{P2}^N) = 0$.

Experimental Results

Experiments Description

We implemented the proposed compilation scheme and have evaluating it by two sets of experiments with different purposes. On the one hand we evaluated the scheme in a satisficing planning context in which we focused on the search for sub-optimal plans using different planning systems, while in the other we focus on the search of optimal plans using admissible heuristics.

Regarding the comparison in the context of the satisficing planning we have considered the following STRIPS+ planning system LAMA[10], Mercury [7], MIPlan [8], IBaCoP2 [3], which are some of the best performing planning system in IPC8 [11], and Fast Downward Stone Soup 2018[?], Fast Downward Remix[?] (abbreviated with FDRemix), which are some of the best performing planning system in the last IPC9 [1] which have been compared with LPRPG-P [4], which is one of the performing planner which supports PDDL3 preferences. Moreover have considered our specifically enhanced version of LAMA for planning with soft

goal, which is LAMA$_P$(h_R), which makes use of admissible heuristic h_R to test the reachability of the soft goals of the problem [9].

As benchmark we have considered the five domains of the qualitative preference track of IPC5 [6] which involve always, sometime, sometime-before, at-most-once and soft goal preferences, i.e Rovers, TPP, Trucks, Openstacks and Storage.

For each original problem all preferences and each original utility were kept. The the classical planners were runned on the compiled problems while LPRPG-P was runned on the original problems of the competition. All the experiments were conducted on a 2.00GHz Core Intel(R) Xeon(R) CPU E5-2620 machine with CPU-time and memory limits of 30 minutes and 8GiB, respectively, for each run of every tested planner. We have tested 8 planners for 5 domains each of which consists of 20 instances for a total of 800 runs.

Table 1 shows the performances of the considered planning system in term of plans quality. As quality measure we have used the IPC quality score, a popular metric of which we have reported a brief description [?].

Given a planner p and a task i we assign, if p solves i, the following score to p:

$$score(p, i) = \frac{cost_{best}(i)}{cost(p, i)}$$

where $cost_{best}(i)$ is the cost of the best know solution for the task i found by any planner, and $cost(p, i)$ is the cost of the solution found by the considered planner p in 30 minutes. In our case our reference for $cost_{best}(i)$ is equal to the cost of the best solution among the tested planners within 30 minutes. If p did not find a solution within the time assigned, then $score(p, i)$ is equal to 0 in order to reward both quality and coverage.

The quality score assigned to each tested planner p is set equal to sum of the quality scores assigned to p over all the considered instances:

$$score(p) = \sum_{i \in tasks} score(p, i)$$

Table 1 reports the quality comparison using the IPC score described above. It is splitted into six parts, at top we have reported the qualitative comparision considering all kinds of preferences together in the computation of the IPC score, while in the remaining subtables we have splitted the have considered each considered class of preference seperately (which is indicated in the header of each subtable).

NOTA (FP): selezionare un numero di pianificatori significativi da includere nella tabella per alleggerire la tabella.

Figures 1a—1e and ??—2a show the qualitative comparison in a different and detailed way. In these histograms we have reported, for each istance of each domain, another quality measure denoted with α_{cost}. Each figure is associated with one of the considered domains and involve two planners, the best performing planner, in term of IPC score and according to Table 1, among the classical planner and LPRPG-P which is the competitor.

We reported a brief description of the metric α_{cost} that we used. Given a planner p and a task i we assign, if p solves i,

the following score to p:

$$\alpha_{cost}(p,i) = cost(p,i)/cost_{total}(i) = \frac{\sum_{P \in \mathscr{P}(i) \,:\, \pi \not\models P} c(P)}{\sum_{P \in \mathscr{P}(i)} c(P)}$$

where $cost(p,i)$ is the cost of the solution found by planner p for the task i within 30 minutes and $cost_{total}(i)$ is the sum of the costs of all the preferences involved in the task i (note that $\mathscr{P}(i)$ denote the set of the preferences of the task i).

If we want to restric the calculation of α_{cost} to a single type of preference, for example just always preferences, we denote the cost as $\alpha_{cost}(\mathcal{A})$ while if nothing is indicated, it means that we have considered all the classes of preferences.

From the previous definition, $\alpha_{cost}(p,i)$ could vary between 0 and 1. If $\alpha_{cost}(p,i) = 0$, then it means that the numerator $cost(p,i)$ is equal to 0 and that p has found an optimal plan for i which satisfies all the preferences of the problem. On the contrary, if $\alpha_{cost}(p,i) = 1$, then it means that p has found the worst plan for i where all the preferences of the problem are violated.

More generally given an instance i, the ratio $\alpha_{cost}(p,i)$, comparing plans produced by different systems, tell us which planner has achieved the satisfaction of the most useful subset of preferences in absolute terms. In particular, the planner with the lowest ratio is the planner who got the best performance on that particular instance.

Satisficing Planning Results

The results obtained comparing the planners considered above show that the compilative approach is almost always preferable since the tested classical planners obtain an higher IPC score than LPRPG-P except for Mercury and MI-Plan.

With reference to Table 1 the compilative approach seems at glance to be particularly preferable in Rovers, Trucks and Storage. In these domains each classical planner performs better or at least comparable than LPRPG-P (except for Mercury in Trucks); IBaCoP2 performs particularly well in Rovers, FDRemix in Trucks and LAMA$_P$(h_R) in Storage. Also MIPlan works well in Trucks but it is penalized due to coverage (it solves only 15 instances out of 20).

The planning system from the more recent IPC9 FDRemix and Fast Downward Stone Soup 2018 perform overall better in this benchmark than the planning system from the previous IPC8 but the enhanced version LAMA$_P$(h_R) is better than everyone else indeed it improves the performance of LAMA in all the considered domains except in Rovers (where there is no soft goal).

Figures 1a–1e show the α_{cost} comparison for each domain comparing the best perfoming classical planner in term of IPC score in such domain with LPRPG-P.

Looking at Figure 1a, where we have reported the α_{cost} comparison for the best performing planners in Rovers in term of IPC score, we can say that the compilative approach, combined with the use of IBaCoP2, achieves to satisfy a better subset of preferences in all the compared instances compared to LPRPG-P.

In TPP the compilative approach seems to be very ineffective, each classical planner achieves an extremely lower quality performance compared to LPRPG-P. The bad performances in this domain are probably due to the many soft goals and sometime preferences because, as shown in [9], the compilation of soft goals can be sometime problematic and neither the use of the reachability heuristic h_R in LAMA$_P$(h_R) can compensate this weakness.

Indeed the part of Table 1, which concerns soft goals and sometime preferences, cleary shows that LPRPG-P is overall more performing than the classical planners especially in TPP in term of these kinds of preferences achieved. This aspect is clearly observable Figures 2a and 2b which report the $\alpha_{cost}(\mathcal{SG})$ and $\alpha_{cost}(\mathcal{ST})$ comparison respectively between LAMA$_P$(h_R) and LPRPG-P in TPP. In these figures we can see that the compilative approach fails to achieve a better subset of soft goals and sometime preferences compared to what LPRPG-P in half of the instances does. But looking at Table 1 we can also observe that LAMA$_P$(h_R) achieves a better result of always, sometime-before and at-most-once preferences compared to LPRPG-P in term of IPC score, but this is not significant because apparently it happens at the expense of soft-goal and sometime preferences which are clearly more expensive to violate (or equivalently more useful to satisfy).

Looking at Figure 1c, where we have reported the α_{cost} comparison for the best performing planners in Trucks in term of IPC score, we can say that the compilative approach, combined with the use of FDRemix, achieves to satisfy a better, worst and equal subset of preferences in all instances compared to LPRPG-P in 15, 4 and 1 respectively.

Regarding Opentacks the tested planners achieve a comparable performance even if the classical planners are slightly penalized compared to LPRPG-P. Looking at Figure 1d, where we have reported the α_{cost} comparison between LAMA$_P$(h_R) and LPRPG-P calculated considering all kind of preferences, we can say that our classical planner achieves to satisfy a better, worst and equal subset of preferenes in 12, 7 and 1 instances respectively

NOTA (FP): scrivi commento Storage.

DOMINO	h^{blind}		h^{max}		$h^{\text{m\&s}}$		h^{cpdb}	
	WRB	Our	WRB	Our	WRB	Our	WRB	Our
Storage	24.78	57.0	29.2	45.0	32.50	24.0	23.10	57.0
Rovers	17.4	24.0	21.43	25.0	16.67	26.0	15.17	23.0
Trucks	18.84	24.0	23.19	25.0	n/a	25.0	n/a	25
TPP	—	47.0	—	45.0	—	47.0	–	40.0

Table 2: Coverage of our and Nebel compilation scheme on the IPC5 benchmarks set with additional instances with random sampled soft-trajectory constraints, A* search for optimal solution. Our results concerning the sampled instances are averaged.

Optimal Planning Results

Similarly to what to what has been done in [12], we have tested our scheme using some admissible heuristics which are h^{blind}, which assign 1 to all states except for goal states to which assign 0, the maximum heuristic h^{max}, the merge and shrink heuristic $h^{\text{M\&S}}$, and the canonical pattern database heuristic h^{cpdb} [?]. These heuristitcs guarantee the optimality of the solution found. Starting from the IPC5 domains, we generated, similar to what was done [12], simpler instances by randomly sampling subsets of the soft trajectory constraints. Starting from each instance we have generated five new instances with 1%, 5%, 10%, 20% and 40% of the (grounded) soft trajectory constraints while the hard goals have remained unchanged.
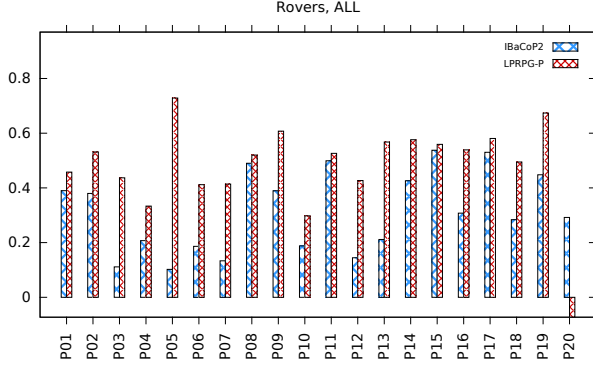
Since we do not have the instances that have been used in the aforementioned paper, we have generated, for each percentage of sampling preferences (except for 100 %), 3 sampled instances in order to average the obtained results and be able to compare with their results. All our experiments were conducted on a 2.00GHz Core Intel(R) Xeon(R) CPU E5-2620 machine with CPU-time and memory limits of 30 minutes and 8GiB, respectively, while the experiments reported in [12] are conducted on a Intel(R) Xeon(R) E5-2650v2 2.60GHz processors with 64GiB with one hour of CPU-time for the search and then our approach is penalized.

The results about this experiment and the comparison with the automata approach are shown Table 2. The results inherent to Openstacks have been excluded because it was not possible to find optimal plans even for the simplest instances.
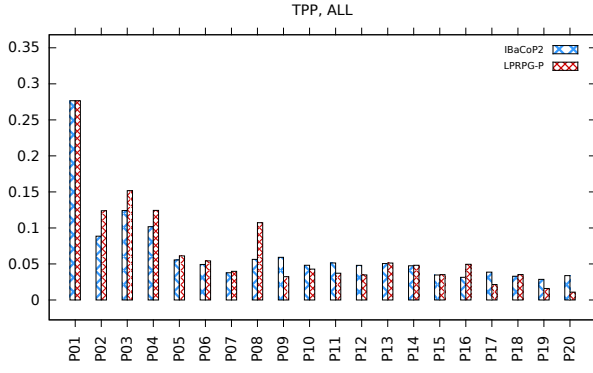
Conclusions

References

[1] The 2018 international planning competition (IPC) classical tracks. https://ipc2018-classical.bitbucket.io. Accessed: 2018-10-12.

[2] M. Briel, R. Sanchez, M. Do, and S. Kambhampati. Effective approaches for partial satisfaction (over-subscription) planning. In *Proceedings of Nineteenth National Conference on Artificial Intelligence (AAAI04)*, pages 562–569, 2004.

[3] I. Cenamor, T. De La Rosa, and F. Fernández. IBaCoP and IBaCoP planner. In *In Eighth International Plan-*

$\mathscr{P}_{\mathcal{ALL}}$

Planner	Rovers	TPP	Trucks	Openstacks	Storage	TOTAL
LAMA$_B$(h_r)	16.98	8.34	15.32	19.28	**18.47**	**78.39**
FDRemix	17.89	7.1	**17.67**	18.99	16.2	77.86
FDSS 2018	17.6	7.03	17.08	18.7	17.11	77.52
LAMA(2011)	17.01	7.53	13.04	18.42	17.81	73.82
IBaCoP2	**19.62**	9.68	10.0	17.85	15.72	72.87
LAMA(2018)	16.44	7.63	13.34	16.03	17.78	71.22
LPRPG-P	11.36	**18.74**	6.99	**19.71**	12.87	69.66
MIPlan	17.65	8.8	9.23	17.35	14.42	67.45
Mercury	16.07	6.57	7.78	18.06	14.5	62.97

$\mathscr{P}_{\mathcal{A}}$

Planner	Rovers	TPP	Trucks	Openstacks	Storage	TOTAL
LAMA$_B$(h_r)	13.46	20.0	15.0	20.0	19.0	87.46
FDSS 2018	13.51	17.0	18.0	17.83	20.0	86.34
LAMA(2011)	13.8	20.0	13.0	20.0	19.0	85.8
IBaCoP2	14.73	20.0	13.0	19.0	19.0	85.73
MIPlan	14.03	20.0	12.0	19.0	20.0	85.03
LAMA(2018)	15.64	20.0	10.0	15.0	19.0	79.64
FDRemix	11.96	15.0	15.0	18.5	19.0	79.46
Mercury	13.07	20.0	4.0	20.0	20.0	77.07
LPRPG-P	13.78	7.0	0.0	19.5	11.0	51.28

$\mathscr{P}_{\mathcal{SG}}$

Planner	Rovers	TPP	Trucks	Openstacks	Storage	TOTAL
LPRPG-P	—	19.45	16.48	19.43	14.94	70.3
FDRemix	—	14.73	17.12	18.63	18.57	69.05
FDSS 2018	—	14.66	16.49	18.4	18.44	67.99
LAMA$_B$(h_r)	—	14.82	14.36	18.7	18.9	66.78
LAMA(2011)	—	13.95	13.57	17.78	18.29	63.6
IBaCoP2	—	16.03	10.7	17.23	18.61	62.57
LAMA(2018)	—	14.82	13.17	15.84	18.3	62.13
MIPlan	—	15.08	9.85	16.7	19.3	60.92
Mercury	—	14.83	7.78	17.39	18.29	58.29

$\mathscr{P}_{\mathcal{AO}}$

Planner	Rovers	TPP	Trucks	Openstacks	Storage	TOTAL
Mercury	16.61	20.0	19.0	—	20.0	75.61
FDSS 2018	16.4	17.0	20.0	—	19.0	72.4
LAMA(2018)	14.92	17.0	20.0	—	20.0	71.92
LAMA$_B$(h_r)	14.54	18.0	20.0	—	19.0	71.54
LAMA(2011)	14.36	17.0	20.0	—	20.0	71.36
FDRemix	16.37	16.0	20.0	—	18.0	70.37
MIPlan	14.03	16.0	15.0	—	20.0	65.03
IBaCoP2	13.21	15.0	16.0	—	19.0	63.21
LPRPG-P	13.42	1.0	19.0	—	12.0	45.42

$\mathscr{P}_{\mathcal{SB}}$

Planner	Rovers	TPP	Trucks	Openstacks	Storage	TOTAL
LAMA$_B$(h_r)	18.25	20.0	18.0	—	19.0	75.25
LAMA(2011)	18.18	20.0	15.5	—	18.0	71.68
Mercury	17.07	20.0	14.5	—	20.0	71.57
MIPlan	18.24	20.0	12.0	—	20.0	70.24
FDRemix	17.49	16.0	16.5	—	20.0	69.99
FDSS 2018	17.46	17.0	16.5	—	19.0	69.96
IBaCoP2	19.72	20.0	12.0	—	18.0	69.72
LAMA(2018)	17.21	20.0	13.33	—	17.0	67.54
LPRPG-P	8.21	14.0	15.5	—	7.0	44.71

$\mathscr{P}_{\mathcal{ST}}$

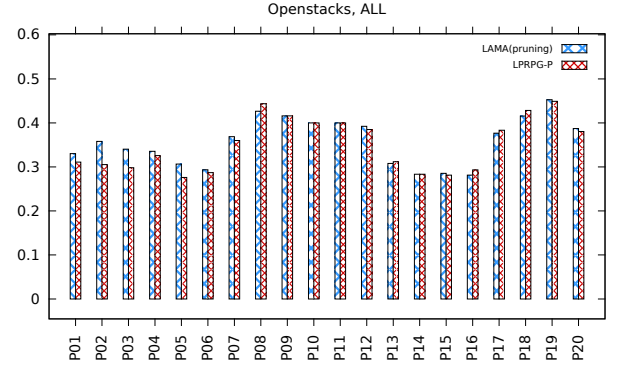| Planner	Rovers	TPP	Trucks	Openstacks	Storage	TOTAL
LAMA(2018)	13.23	11.0	—	—	20.0	44.23
LAMA$_B$(h_r)	14.28	10.0	—	—	19.0	43.28
FDSS 2018	16.17	8.0	—	—	19.0	43.17
IBaCoP2	16.81	10.0	—	—	16.0	42.81
LAMA(2011)	14.41	9.0	—	—	19.0	42.41
LPRPG-P	9.56	17.0	—	—	14.0	40.56
FDRemix	15.44	8.0	—	—	16.0	39.44
MIPlan	14.91	9.0	—	—	14.0	37.91
Mercury	12.78	4.0	—	—	12.0	28.78

Table 1: Temp caption

(a) Rovers

(b) TPP

(c) Trucks

(d) Openstacks

(e) Storage

Figure 1: Quality Comparison using α_{cost} for each domain. Each bar represents the α_{cost} of the best plan produced by the considered planner. The negative bar represents an instance which has not been solved or that has no preferences of that kind. From the top to bottom we have provided the results about α_{cost} calculated considering each kind of preferences for Rovers, TPP, Trucks, Openstacks and Storage.

ning Competition Booklet (ICAPS-14), pages 35–38, 2014.

[4] Amanda Jane Coles and Andrew Coles. Lprpg-p: Relaxed plan heuristics for planning with preferences. In *Proceedings of the Twenty-First International Conference on Automated Planning and Scheduling (ICAPS 2011)*, pages 26–33, 2011.

[5] M.B. Do and S. Kambhampati. Partial satisfaction (over-subscription) planning as heuristic search. In *Proceedings of Fifth International Conference on Knowledge Based Computer Systems (KBCS04)*, page mancanti, 2004.

[6] A. Gerevini, P. Haslum, D. Long, A. Saetti, and Y. Dimopoulos. Deterministic planning in the fifth international planning competition: PDDL3 and experimental evaluation of the planners. *Artificial Intelligence*, 173(5-6):619–668, 2009.

[7] M. Katz and J. Hoffmann. Mercury planner: Push-
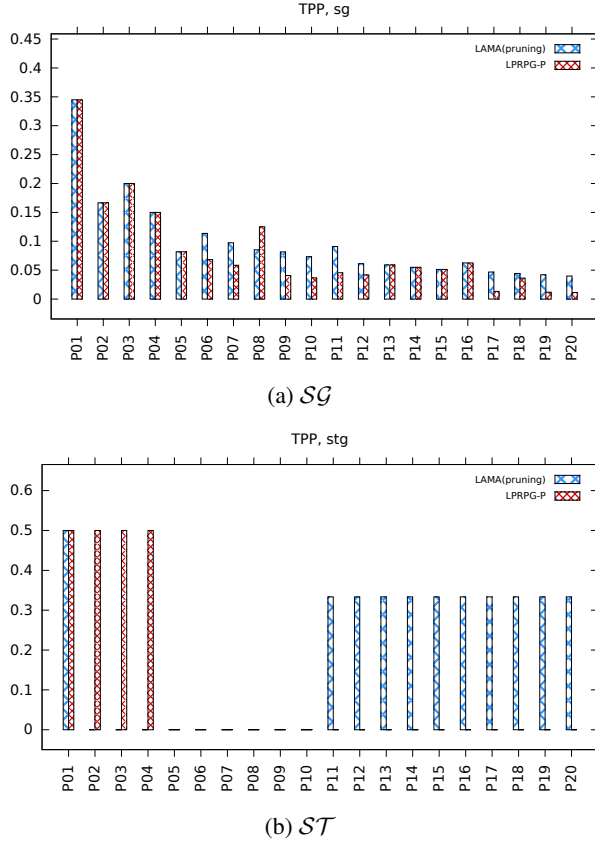
(a) $\mathcal{SG}$

(b) $\mathcal{ST}$

Figure 2: Each bar represents the α_{cost} of the best plan produced by the considered planner. The negative bar represents an instance which has not been solved or that has no preferences of that kind. From left to right we have provided the results about the α_{cost} calculated considering each kind of preferences, always, at-end (or soft goals), at-most-once, sometime-before and sometime preferences.

ing the limits of partial delete relaxation. In *In Eighth International Planning Competition Booklet (ICAPS-14)*, pages 43–47, 2014.

[8] S. Núñez, D. Borrajo, and C. Linares López. MIPlan and DPMPlan. In *In Eighth International Planning Competition Booklet (ICAPS-14)*, pages 13–16, 2014.

[9] F. Percassi, A. Gerevini, and H. Geffner. Improving plan quality through heuristics for guiding and pruning the search: A study using LAMA. In *Proceedings of the Tenth International Symposium on Combinatorial Search (SoCS 2017)*, pages 144–148, 2017.

[10] S. Richter and M. Westphal. The LAMA planner: Guiding cost-based anytime planning with landmarks. *Journal of Artificial Intelligence Research*, 39(1), 2010.

[11] M. Vallati, L. Chrpa, M. Grześ, T. L. McCluskey, M. Roberts, and S. Sanner. The 2014 international planning competition: Progress and trends. *AI Magazine*, 36(3):90–98, 2015.

[12] Mattmller R. Wright B. and Nebel B.. Compiling away soft trajectory constraints in planning. In *In Proceedings of the Sixteenth Conference on Principles of Knowledge Representation and Reasoning (KR18)*, 2018.